

AI 行业周报

多源聚合 · 真实数据 · 情感分析 · 趋势洞察 · 数据看板

2026-06-15 ~ 2026-06-21

共 88 条新闻 · 10 个数据源

 KPI 概览

88 本周新闻	10 数据来源	59 利好信号	5 利空信号	2 中国动态
------------	------------	------------	-----------	-----------

 情感分析

59 利好/偏多 (67%)	24 中性 (27%)	5 利空/偏空 (6%)
-------------------	----------------	-----------------

偏多

行业首款开源鸿蒙消费级人形机器人：松延动力 N2 亮相华为开发者大会 HDC 2026

IT之家 6 月 15 日消息，行业首款开源鸿蒙消费级人形机器人——松延动力的 OpenHarmony 版 N2 亮相上周末的华为开发者大会 HDC 2026。松延动力表示，当前消费级人形机器人普遍面临应用生态与长期发展的多重痛点。多数产品仍停留在独立硬件终端的形态，难以与手机、智慧屏、智能家居、教育硬件等主流智能设备实现深度互联互通，久而久之形成生态孤岛，极大限制了多场景融合体验的打造。与此同时，传统系统架构下的跨设备开发流程繁琐，网络协议适配、异构数据转换等工作大幅抬高研发门槛与成本，拉长产品迭代周期。再加之外置系统普遍体积偏大、资源占用率高，与人形机器人轻量化、低功耗的行业发展趋势相悖，这些问题都成为制约消费级人形机器人规模化普及、业态升级的关键阻碍。松延动力产品副总裁唐子洋认为，依托开源鸿蒙微内核架构、高安全性与跨终端联动优势，AI 系统智能体语音交互控制、长程任务自规划能力，机器人未来可与家庭场景下的各类智能终端无缝协作，既能化身家庭智能管家，联动家电、安防设备打造全屋智能体验，也能依托多设备协同能力，强化教育陪伴、互动教学等核心场景体验。IT之家从大会获悉，松延动力还宣

IT之家 · 2026-06-15T03:56

偏多

东风纳米 01 智趣版纯电小车 6 月 16 日上市，CLTC 纯电续航 305km

IT之家 6 月 15 日消息，东风汽车宣布旗下纳米 01 纯电小车智趣版将于 6 月 16 日（明天）正式上市，该车 CLTC 续航里程为 305km，主要在内饰方面同步新近上市的纳米 01 Cross 设计，换用了长条形的全液晶仪表盘以及尺寸更大的中控大屏。外观方面，该车前脸配备分布式灯组，拥有贯穿式日间灯；车身使用隐藏式门把手和低风阻轮毂；汽车尾灯棱角分明，内置圆形 LED 光源，高位扰流板经过上扬处理，不乏运动感。该车座舱采用了纳米 01 Cross 同款设计，换用 15.6 英寸 2.5K 高色域超清护眼智能屏、8.88 英寸液晶组合仪表，有望支持 540° 全景鸟瞰、全场景泊车与循迹倒车等 L2 级组合辅助驾驶功能。规格方面，该车长宽高分别为 4020x1810x1570mm，轴距为 2663mm。不过官方尚未公布新款车型的动力信息，作为参考，现款在售车型采用的单电机前驱布局，电机最大功率为 70kW，峰值扭矩为 160N·m。悬架系统采用了该级别车型中常见的前麦弗逊后扭力梁的组合。

IT之家 · 2026-06-15T03:26

偏多

小米米家扫拖机器人 6 Pro 国补探新低：水箱版 2192 元、上下水版 2518 元

小米米家扫拖机器人 6 Pro 开售于 2026 年 4 月 8 日，水箱版定价 3799 元、上下水版定价 4199 元。京东 618 大促期间，按下方步骤下单，叠加 PLUS 立减 + 15% 国补后，水箱版 2192.34 元、上下水版 2518.72 元：满 1 元立打 8.2 折券：点此查看。加入购物车时选择“以旧换新”补贴 200 元（可选电动牙刷），京东回收不看旧物品相，老、旧、坏电器皆可回收。晒单再返 20 元京豆。京东小米 米家扫拖机器人 6 Pro 水箱版券后 2192.34 元领 8.7 折券京东小米 米家扫拖机器人 6 Pro 上下水版券后 2518.72 元领 8.7 折券京东无门槛红包至高 26618 元，支持即领即用：点此抽取。IT之家从商品页面获悉，这款新品支持 AI 三摄全景识别，可识别数据线、拖鞋、沙发等 280 种常见障碍物和家具，最细可识别 3mm 耳机线；自研 AI 脏污识别算法，高效识别猫砂、狗粮、牛奶等 47 种干湿脏污，并智能匹配清洁策略；配备高清 AI 摄像头，可用米家 App 在视频模式下，随时查看宠物状态。新品支持小米澎湃

IT之家 · 2026-06-15T02:57

偏多

山灵 Moria「生命之树」钛金振膜单动圈耳机本周上市：搭复合振膜单元、钛冰晶工艺面板，4998 元

IT之家 6 月 15 日消息，山灵现已公布一款 Moria「生命之树」钛金振膜单动圈耳机，将于本周上市，定价为 4998 元。该耳机面板采用独特的“钛冰晶”工艺，每片纹理均独一无二，视觉效果如极夜星河或冰面裂纹，具有一定艺术感。同时耳机腔体由高精度 CNC 纯钛打造，轻盈坚固。其标配 5N 长结晶无氧铜镀银线，支持 0.78mm 2PIN 接口，方便用户更换升级线材。官方强调，该耳机搭载“超级动圈”单元，核心为一枚经过 60 层溅射工艺打造的钛金复合振膜，结合 PEK 基材与物理真空直流磁控溅镀技术，在保持柔顺性的同时提升刚性，显著增强解析力。同时，该耳机采用三内腔双磁路动圈架构，通过两个独立后腔室优化气流控制，配合双钕磁铁系统，号称能够实现均衡三频响应与自然的声场动态。

IT之家 · 2026-06-15T02:56

偏多

OPPO ColorOS 16 六月更新公布：锁屏岛新增酷狗音乐等，6 月 30 日前完成推送

IT之家 6 月 15 日消息，OPPO 今日公布 ColorOS 16 六月更新内容：AI 流体云再上新，锁屏岛新增酷狗音乐，流体云新增超级课程表、日历课程提醒；小布建议足球观赛卡，小布助手观赛智能体；还有飞牛私有云无感连接、OPPO Bubble 潮玩自拍屏、全新端午节水印等更多功能。IT之家整理如下：AI 流体云新增支持在小布建议和流体云上，显示超级课程表应用的课程信息支持在小布建议和流体云上，显示支付宝停车服务的计费变动支持在流体云 and 锁屏岛上，显示酷狗音乐歌词AI 功能新增小布建议足球观赛卡支持通过内容传送门，将图片、文本、文档拖拽存入小布记忆小布助手接入京东、京东金融、飞常准、移动咪咕智能体服务，支持展示对应服务卡片，点击即可查看详情无需跳转第三方应用支持在小布建议上，显示天气预警服务小布建议支持在桌面添加 AI 简报卡片优化小布记忆旅行合集功能，旅行地图支持依据行程生成路线规划图，并在地图中同步呈现景点周边美食小布助手的圈选体验，支持调整圈选框大小与位置，升级圈选光效通信与互联新增设备空间支持添加和管理飞牛私有云设备（NAS）相册支持与飞牛私有云（NAS）无感连接，并可在系

IT之家 · 2026-06-15T02:48

利好

长城哈弗 H6 经典版 2026 款上市，限时优惠换新价 8.19 万元起

IT之家 6 月 15 日消息，长城哈弗 H6 经典版 2026 款今日正式上市，上市价 9.99 万元，限时优惠换新价 8.19 万元起。该车限时可享 8000 元现金礼、最高 1 万元置换补贴，以及首任个人非营运车主发动机、变速器的核心零部件终身质保。长城哈弗 H6 经典版 2026 款新增“远山灰”外观配色，尺寸为 4653*1886*1730mm，轴距 2738mm；配有 1.5T 发动机 + 7DCT 湿式双离合变速器，最大功率 135kW，最大扭矩 275N·m。该车配有 14.6 英寸中控屏、

手车互联，中控 UI 升级、高德地图红绿灯倒计时；支持 ACC 自适应巡航、540° 全景影像。IT之家查询配置表获悉，长城哈弗 H6 经典版 2026 款搭载前麦弗逊式独立悬架 / 后多连杆式独立悬架，标配 6 安全气囊，车身高强钢占比约 72%。该车还新增电尾门、前排座椅电动调节、前排座椅加热，配有可开启全景天窗、8 扬声器、外后视镜加热以及自动折叠。

IT之家 · 2026-06-15T02:46

模型发布 偏多

Claude Code Guide 2026: 25 Features with Examples + Demo

Claude Code is a layered agentic coding tool, not a single chat prompt. This guide breaks down 25 features, from CLAUDE.md, skills, subagents, and hooks to MCP and Auto Mode. It includes a comparison table, working code examples, real use cases, and an interactive demo you can try. The post Claude Code Guide 2026: 25 Features with Examples + Demo appeared first on MarkTechPost.

MarkTechPost · 2026-06-15T01:04

监管政策 偏多

A Deep Reinforcement Learning (DRL)-Based Transformer Method for Solving the Open Shop Scheduling Problem

arXiv:2606.13682v1 Announce Type: new Abstract: The open shop scheduling problem (OSSP) arises in many industrial and service settings but remains computationally challenging as the number of jobs and machines increases. While exact methods quickly become intractable, classical dispatching rules and metaheuristics may require substantial tuning to maintain solution quality at large scales. This study develops a Transformer-based scheduling policy for OSSP using an encoder-decoder architecture w

arXiv AI · 2026-06-15T00:00

模型发布 AI Agent 偏多

Orchestra-o1: Omnimodal Agent Orchestration

arXiv:2606.13707v1 Announce Type: new Abstract: The recent success of agent swarms has shifted the paradigm of large language model (LLM)-based agents from single-agent workflows to multi-agent systems, highlighting the importance of agent orchestration for task decomposition and collaboration. However, existing orchestration frameworks are limited to a narrow set of modalities and struggle to generalize to more complex settings where heterogeneous modalities coexist and interact. This limitati

arXiv AI · 2026-06-15T00:00

AI Agent 偏多

Hybrid Open-Ended Tri-Evolution Makes Better Deep Researcher

arXiv:2606.13710v1 Announce Type: new Abstract: Deep research and agent evolution serve as de-facto tasks for AI agents in real-world applications toward artificial general intelligence. The former enables autonomous retrieval and integration of information in open-ended environments to tackle open-ended research tasks, yet it is constrained by the static parametric deep research capabilities of agent systems. The latter allows agents to autonomously interact with the environment to gain experi

arXiv AI · 2026-06-15T00:00

模型发布 AI Agent 开源/评测 偏多

WorkBench Revisited: Workplace Agents Two Years On

arXiv:2606.13715v1 Announce Type: new Abstract: The best agent on WorkBench in March 2024, GPT-4, completed 43% of tasks and took an unintended harmful action, such as emailing the wrong person, on 26% of them. We re-visit the benchmark in June 2026 and find that the best agent to date, Claude Opus 4.8, completes 89% and takes an unintended harmful action on 2.5%. Aside from this considerable progress in frontier agent performance, three things stand out. First, capability and safety go togethe

arXiv AI · 2026-06-15T00:00

偏空

Refusal Beyond a Single Direction: A Preliminary Comparison of Diff-in-Means and INLP

arXiv:2606.13720v1 Announce Type: new Abstract: Arditi et al. (2024) has shown that refusal in safety fine-tuned chat models is mediated by a single linear direction in the residual stream, recoverable by a difference-in-means (DiM) of harmful and harmless activations. We compare DiM-based interventions (activation addition and directional ablation) with two interventions derived from Iterative Nullspace Projection (INLP) -- nullspace projection and counterfactual flipping -- on five open-weigh

arXiv AI · 2026-06-15T00:00

模型发布 AI Agent 偏多

YeasierAgent: Agentic Social Sandbox as a Canvas for Intent-Driven Creation of Platform-Agnostic Symbiotic Agent-Native Applications

arXiv:2606.13722v1 Announce Type: new Abstract: This paper introduces YeasierAgent, an application-building paradigm based on symbiotic agents, narrative worlds, and scene-aware interaction. It challenges the conventional device-coupled model of software by redefining

applications as collaborative spaces among users, agents, and worlds. We present a system architecture that achieves two primary contributions: (1) enabling the rapid, cross-platform construction of agent-native applications by ut

arXiv AI · 2026-06-15T00:00

偏多

TwinBI: An Agentic Digital Twin for Efficient Augmented Interactions with Business Intelligence Dashboards

arXiv:2606.13731v1 Announce Type: new Abstract: Business intelligence (BI) increasingly combines dashboard interaction with LLM-based assistance, but these two modes often fall out of sync during multi-step analysis. As users switch between direct dashboard manipulation and natural-language queries, it becomes difficult to preserve a consistent analytical state across filters, hierarchies, metrics, and chart context. We present TwinBI, an agentic digital-twin framework that couples an LLM-based

arXiv AI · 2026-06-15T00:00

偏多

AI Receptivity or AI Adoption Breadth? A Tool-Specific Reanalysis of the Lower-Literacy/Higher-Usage Link

arXiv:2606.13734v1 Announce Type: new Abstract: Recent evidence reported by Tully, Longoni, and Appel (2025) suggests that lower artificial intelligence (AI) literacy predicts greater receptivity toward AI. We revisit this claim using the public data from Study 3 of that article, which measures past usage of five AI tool categories on a five-point frequency scale. We first reproduce the negative association between AI literacy and aggregate AI usage using OLS on participant-level averages, bina

arXiv AI · 2026-06-15T00:00

模型发布 监管政策 偏多

Capability Minimization as a Safety Primitive: Risk-Aware Causal Gating for Least-Privilege LLM Agents

arXiv:2606.13884v1 Announce Type: new Abstract: Modern decision systems increasingly rely on learned components whose outputs may be confident yet wrong, exposing downstream actions to costly errors. We introduce Risk-Aware Causal Gating (RACG), a framework that decides whether to act on, defer, or abstain from a model's prediction by combining causal effect estimation with calibrated risk control. RACG models the causal pathway from candidate actions to outcomes and gates each decision accordi

arXiv AI · 2026-06-15T00:00

AI Agent 中性

A Multi-Agent AI System for Automated High School Transcript Processing: Collaborative Document Analysis at Scale

arXiv:2606.13916v1 Announce Type: new Abstract: Each year, college admissions offices face an overwhelming challenge: processing millions of high school transcripts, each with unique formats, grading systems, and layouts. This manual process creates operational bottlenecks that delay admissions decisions and consume valuable resources. We present a transformative solution through a multi-agent AI system where specialized agents collaborate to automatically process diverse transcript formats thr

arXiv AI · 2026-06-15T00:00

模型发布 中性

Adversarial Concept Search: Predicting Compositional Errors From Feature Geometry

arXiv:2606.13934v1 Announce Type: new Abstract: Humans cannot always intuit what scenarios are most challenging to LLMs. Hoping to capture challenging edge cases, developers either design problems to be difficult for humans or curate extensive benchmarks. What if we could instead anticipate which scenarios a model will fail on? In this paper, we use an LLM's representational geometry to predict which concept combinations it will fail on. We attribute this compositional failure to interference b

arXiv AI · 2026-06-15T00:00

AI Agent 中性

Minim: Privacy-Aware Minimal View for Agents via Trusted Local Sanitization

arXiv:2606.13949v1 Announce Type: new Abstract: Modern LLM-powered autonomous agents increasingly rely on rich user interface (UI) state observations to achieve reliable action grounding in complex digital environments. However, many deployments transmit the full UI state to remote inference servers even when most elements are irrelevant to the current task, which can leak sensitive but unnecessary context such as authentication codes, private notifications, and background application states. W

arXiv AI · 2026-06-15T00:00

AI Agent 偏多

Formalizing Numerical Analysis: An Agent Pipeline and Quality Audit Beyond Kernel Acceptance

arXiv:2606.14000v1 Announce Type: new Abstract: Recent work has demonstrated that coding agents can formalize entire advanced mathematics textbooks in Lean 4, yet existing efforts concentrate on branches of mathematics already well-represented in mathlib and measure success solely

through kernel acceptance. We address both limitations by applying a coding agent to formalize Numerical Methods for Ordinary Differential Equations, a textbook in numerical analysis that is largely absent from mathli

arXiv AI · 2026-06-15T00:00

中国动态 偏空

China may have accessed Mythos

According to a new report from Semafor, the White House's decision to impose export restrictions on Anthropic's Mythos was driven in part by fears that it had been accessed by a group linked to China. If the Chinese government actually had access to Mythos 5 or Fable 5, it would present a serious national security [~]

The Verge · 2026-06-14T14:27

模型发布 偏多

Databricks Open-Sources Omnigent: A Meta-Harness That Composes, Governs, and Shares AI Agents Across Claude Code, Codex, and Pi

Databricks has open-sourced Omnigent, a meta-harness that sits above coding agents like Claude Code, Codex, and Pi. It adds composition, contextual policies, and live session sharing under one interface, on terminal, web, desktop, and mobile. The Apache 2.0 project is in alpha. The post Databricks Open-Sources Omnigent: A Meta-Harness That Composes, Governs, and Shares AI Agents Across Claude Code, Codex, and Pi appeared first on MarkTechPost.

MarkTechPost · 2026-06-14T05:01

偏空

Amazon security research reportedly led to the White House' s Anthropic Fable ban

According to the Wall Street Journal, the export control directive that led to Anthropic cutting off access to Fable 5 and Mythos 5 was triggered in part by cybersecurity research from Amazon and conversations between CEO Andy Jassy and the White House. According to the report, the paper from Amazon claims that, through a series [~]

The Verge · 2026-06-13T17:39

模型发布 AI Agent 偏多

How to Build a QwenPaw Agent Workspace with Custom Skills, Model Providers, Console Access, and Streaming API Testing

In this tutorial, we implement a QwenPaw workflow that provides a practical environment for building and testing an agent-powered assistant. We install and initialize QwenPaw, configure its working directory, set up authentication, connect optional model providers via Colab secrets, and create a structured workspace with custom skills and local knowledge files. We also launch the [~] The post How to Build a QwenPaw Agent Workspace with Custom Skills, Model Providers, Console Access, and St

MarkTechPost · 2026-06-13T17:27

模型发布 中性

My yard is dying, so I made an app for that

When I returned to my computer five minutes after giving Gemini a lengthy prompt, I had two things: a functional app in a preview window, and a message about a bug. "~ Channel is unrecoverably broken and will be disposed!" Sounded bad! But right below it was a button to fix the bug. Pretty weird [~]

The Verge · 2026-06-13T09:00

中性

A Court Has Ruled That Google Is Liable for False Statements Generated by AI Overviews

The ruling holds that a company that designs, trains, operates, and manages an AI system must assume legal liability for any damages caused by the responses it generates.

Wired · 2026-06-13T09:00

中性

Anthropic cuts off Fable 5 and Mythos 5 access following government order

On Friday evening, the government ordered Anthropic to block access to Fable 5 and Mythos 5 for all foreign nations, both inside and outside the US, due to national security concerns. That order included employees of Anthropic. To meet those demands, the company has completely cut off access to the models for all customers. In [~]

The Verge · 2026-06-13T08:55

模型发布 偏多

Anthropic Disables Claude Fable 5 and Mythos 5 After US Government Order

shutdown followed a US government export control directive citing national security authorities. All other Anthropic models, including Opus 4.8, remain available. The post Anthropic Disables Claude Fable 5 and Mythos 5 After US Government Order appeared first on MarkTechPost.

MarkTechPost · 2026-06-13T08:15

偏多

Apple’ s new AI photo editing tools mostly work, for better and worse

The most popular camera in the world just got its first set of serious AI photo editing features, and I don't think any of us are ready. As far as AI photo editing goes, the new features in iOS 27 are pretty tame compared to what you can do on, say, Google's Pixel phones. But [8230;]

The Verge · 2026-06-13T08:00

偏空

The future of Hollywood isn’ t feeding prompts into vanilla gen AI models

For all the noise that's been made about how generative AI is poised to revolutionize the filmmaking industry, there haven't really been any projects created with the technology that felt like the sort of entertainment people would pay to see. Most AI firms' video models are still only capable of churning out short bursts of [8230;]

The Verge · 2026-06-13T07:00

模型发布 利好

Moonshot AI Releases Kimi K2.7-Code: a Coding Model Reporting +21.8% on Kimi Code Bench v2 Over K2.6

Moonshot AI has open-sourced Kimi K2.7-Code under a Modified MIT license. It is a coding-focused, agentic model built on Kimi K2.6, with a 256K context window and roughly 30% lower reasoning-token usage. Moonshot reports gains over K2.6 on six benchmarks, including +21.8% on Kimi Code Bench v2. The model is available via the Kimi API and Kimi Code. The post Moonshot AI Releases Kimi K2.7-Code: a Coding Model Reporting +21.8% on Kimi Code Bench v2 Over K2.6 appeared first on MarkTechPost.

MarkTechPost · 2026-06-13T04:57

模型发布 偏多

Anthropic Says It’ s Taking Claude Fable 5 Offline to Comply With US Government Order

“The government believes it has become aware of a method of bypassing, or ‘jailbreaking’ Fable 5,” the company said in a blog post.

Wired · 2026-06-13T02:26

中性

Meta Employees Absolutely Hate Mark Zuckerberg’ s Plan for a Companywide AI Hackathon

“I’ m not sure that this company supports a hackathon culture anymore,” one employee posted in a forum open to the entire staff.

Wired · 2026-06-12T23:45

中性

‘Tell Him He’ s a Piece of Shit’ : Meta’ s New AI Unit Is a Total Mess

Executives and employees alike are struggling with Meta’ s chaotic AI strategy, according to sources and internal discussions reviewed by WIRED.

Wired · 2026-06-12T21:16

模型发布 利好

Google Releases Gemini-SQL2: Gemini 3.1 Pro Text-to-SQL Scores 80.04% on BIRD Single-Model Leaderboard

We look at Gemini-SQL2, the text-to-SQL capability Google Research announced on June 12, 2026. Powered by Gemini 3.1 Pro, it posted 80.04% execution accuracy on the BIRD single-model leaderboard. We explain what the score measures, how the leaderboard stacks up, and what Google has not yet disclosed. We also cover use cases and a schema-grounded implementation pattern. The post Google Releases Gemini-SQL2: Gemini 3.1 Pro Text-to-SQL Scores 80.04% on BIRD Single-Model Leaderboard appeared first o

MarkTechPost · 2026-06-12T21:04

中国动态 偏多

China Didn’ t Make Americans Hate Data Centers

GOP lawmakers, tech investors, and even OpenAI have tied the anti-data-center movement in the US to Chinese interference. Experts say it’ s much more complicated than that.

Wired · 2026-06-12T17:17

偏多

Coinbase for Agents: Automating portfolio trading with AI

Coinbase for Agents connects AI to financial execution channels to automate trading and payments directly from user portfolios. Large language models process vast quantities of data but lack direct integration with active financial portfolios. Individuals frequently employ these models to evaluate market developments or research investment opportunities. These software tools possess the capacity for complex [8230;] The post Coinbase for Agents: Automating portfolio trading with AI appeared fir

AI News · 2026-06-12T15:34

投融资 利好

SpaceX’ s massive IPO: all the latest news

SpaceX’s IPO on Friday allows the public to buy shares of the combined rocket, AI, and social media company for the first time, and raised enough money to make Elon Musk the first trillionaire. He has more wealth, on paper at least, than the economies of nations like Ireland, Sweden, or his home country of [8230;]

The Verge · 2026-06-12T10:35

偏多

Jeff Bezos’ AI startup aims to build an ‘artificial general engineer’

Amazon founder Jeff Bezos says his new AI startup will work toward developing an "artificial general engineer," according to reports from The New York Times and CNBC. The startup, called Prometheus, aims to develop AI-powered engineering tools to aid in the design of physical products. The NYT first reported on Prometheus last November, but now [8230;]

The Verge · 2026-06-12T10:03

AI Agent 偏多

Moonshot AI Launches Kimi Work, a Local Desktop Agent Reportedly Running on Kimi K2.6 With a 300-Sub-Agent Agent Swarm

Moonshot AI's Kimi Work is a local desktop agent for macOS and Windows. It runs a 300-sub-agent swarm, drives your logged-in browser via WebBridge, and schedules background jobs. The post Moonshot AI Launches Kimi Work, a Local Desktop Agent Reportedly Running on Kimi K2.6 With a 300-Sub-Agent Agent Swarm appeared first on MarkTechPost.

MarkTechPost · 2026-06-12T09:37

模型发布 偏多

Zyphra Release Zamba2-VL: Hybrid Mamba2–Transformer Vision-Language Models That Cut Time-to-First-Token by About an Order of Magnitude

Zyphra has released Zamba2-VL, a family of open vision-language models at 1.2B, 2.7B, and 7B parameters. The models use a hybrid Mamba2 state-space and Transformer backbone, shipping under Apache 2.0. They stay competitive with comparable Transformer VLMs while cutting time-to-first-token by about an order of magnitude. The post Zyphra Release Zamba2-VL: Hybrid Mamba2–Transformer Vision-Language Models That Cut Time-to-First-Token by About an Order of Magnitude appeared first on MarkTechPost.

MarkTechPost · 2026-06-12T08:06

中性

Apple’ s Camera Chief Thinks AI Can Give You Superpowers

The generative features in iOS 27’ s new Photos app will add fake pixels to some of your shots, but Apple’ s Jon McCormack says the company isn’ t using AI “for the sake of AI.”

Wired · 2026-06-12T01:00

投融资 偏多

Why You Might Already Own SpaceX Shares, Siri’ s AI Makeover, and Knicks Owner’ s Surveillance Machine

Today on Uncanny Valley, we take an early look at the SpaceX IPO and why you might find yourself among the investors without even realizing it.

Wired · 2026-06-11T23:37

利好

Meet the OpenAI Engineer Leading ChatGPT’ s Biggest Transformation Yet

Thibault Sottiaux helped make AI coding one of OpenAI’ s fastest-growing businesses. Now he’ s overseeing a sweeping overhaul of ChatGPT.

Wired · 2026-06-11T21:21

AI Agent 基础设施 偏多

Visa ChatGPT integration enables AI agent retail purchasing

Visa has linked its payment infrastructure to ChatGPT, enabling AI agents to recommend retail products and execute financial transactions. The deployment removes human intervention from the final stages of the retail funnel. Autonomous agents will now process user prompts, evaluate merchant catalogues, and complete the checkout process using Visa's payment rails at any supporting merchant. [8230] The post Visa ChatGPT integration enables AI agent retail purchasing appeared first on AI News.

AI News · 2026-06-11T11:08

投融资 偏多

Google DeepMind is worried about what happens when millions of agents start to interact

Google DeepMind is funding research into the potential dangers of situations where millions of different AI agents interact with each other online. According to Rohin Shah, who directs the company's AGI safety and alignment research, the mass-market arrival of agents that can carry out tasks without human oversight and follow instructions given to them by other8230;

MIT科技评论 · 2026-06-11T11:00

偏多

Xebia: Why AI agents fail without the right data foundation

If your remit is to help your organisation add AI agents to accelerate its processes, you have to start at the foundation – and that means making your data available for AI consumption. Agentic AI scales on data strength, as Niels Zeilemaker, global CTO at Xebia, explains. "If you don't think about that, you can [8230] The post Xebia: Why AI agents fail without the right data foundation appeared first on AI News.

AI News · 2026-06-11T10:42

偏多

Siri AI arrives with Google inside, and much of the world is locked out

8220;We8217;ve all had that moment where you search for something you know is there, but it just won8217;t show up.8221; Apple8217;s Stacey Ford, vice president of OS Program Management, was talking about Spotlight at WWDC 2026, but she could have been describing the company8217;s AI ambitions.8160; On Monday at Apple Park, the thing that wouldn8217;t [8230] The post Siri AI arrives with Google inside, and much of the world is locked out appeared first on AI News.

AI News · 2026-06-10T11:00

偏多

McDonald's tests Google-backed AI drive-thru ordering system

McDonald's is testing a new AI system that can take drive-thru orders and support restaurant operations. The system, called ArchIQ and nicknamed "Archy," was introduced during the company's Worldwide convention, according to Restaurant Business. It is being tested at five McDonald's locations in the United States, though the company has not named the restaurants involved. [8230] The post McDonald's tests Google-backed AI drive-thru ordering system appeared first on AI News.

AI News · 2026-06-10T10:00

偏多

How to sign PDFs easily online with a PDF signer

Signing PDFs has become an important task for businesses and individuals alike. Whether you8217;re handling contracts, legal agreements, or forms, the ability to quickly and securely sign PDFs online is essential. Fortunately, with the rise of online PDF signers, signing PDFs has never been easier. Common challenges in signing PDFs Signing PDFs might seem straightforward, [8230] The post How to sign PDFs easily online with a PDF signer appeared first on AI News.

AI News · 2026-06-09T09:01

AI Agent 中性

Autonomous AI Data Loss in DevOps: Building Efficient Defenses

Autonomous AI agents are altering the speed at which software is shipped. Unfortunately, they are also shrinking the time it takes for a mistake to become a catastrophe, creating a dangerous blind spot in many security strategies. The threat no longer comes just from external ransomware or malicious insiders. It comes from authorized, internal tools. [8230] The post Autonomous AI Data Loss in DevOps: Building Efficient Defenses appeared first on AI News.

AI News · 2026-06-09T08:50

偏多

Aviva deploys AI to stop £230M in sophisticated insurance fraud

Aviva has uncovered a record £230 million in insurance fraud claims and is using AI tools to counter the growing problem. The battleground has changed, and the culprits are also coming armed with a new generation of tools. We8217;re now in an environment where AI is being used not just to defend against fraud, but [8230] The post Aviva deploys AI to stop £230M in sophisticated insurance fraud appeared first on AI News.

偏多

Weis Markets adds Instacart AI-powered shopping carts to stores

Weis Markets is adding Instacart’s AI-powered shopping carts, Caper Carts, to select stores in Pennsylvania, bringing digital coupons, loyalty features, and repeat-purchase recommendations into the grocery aisle. The Pennsylvania-based grocery chain is working with Instacart to deploy the smart carts, which include cameras, certified scales, location systems, and a touchscreen. According to Instacart, Caper Carts [8230;] The post Weis Markets adds Instacart AI-powered shopping carts to stores

AI News · 2026-06-08T10:00

偏多

The latest AI news we announced in May 2026

May AI recap

Google AI Blog · 2026-06-05T14:45

偏多

How C3 AI agents will automate predictive maintenance for Shell

Shell will use agents from C3 AI to shift from basic anomaly detection towards fully-automated predictive maintenance. The global energy giant is building on their current use of the C3 AI Reliability Suite, which already keeps tabs on more than 30,000 crucial pieces of equipment across upstream and downstream operations. Shell now intends to lean [8230;] The post How C3 AI agents will automate predictive maintenance for Shell appeared first on AI News.

AI News · 2026-06-05T11:34

AI Agent 偏多

Meta Business Agent drives AI-powered conversational commerce

Meta has launched Business Agent to automate conversational commerce workflows directly inside its messaging applications. The software allows global retail brands to execute transactions and field support tickets without human intervention. Deploying this architecture places agentic AI directly at the core of social commerce. Meta integrated these workflows natively into Instagram, Messenger, and soon WhatsApp. [8230;] The post Meta Business Agent drives AI-powered conversational commerce app

AI News · 2026-06-04T17:17

偏多

Scout from M’ Soft is the agentic Autopilot that works across M365

Microsoft has announced the wider testing of its new Autopilot feature at the Microsoft Build event this week, backed by a post on the company’s website. Autopilots are described as a new category of agents that can work autonomously on a user’s behalf. Microsoft says each Autopilot has its own identity, and so multiple agents [8230;] The post Scout from M’s Soft is the agentic Autopilot that works across M365 appeared first on AI News.

AI News · 2026-06-04T11:41

中性

5 ways Google Search can level up your thrift and vintage shopping

A colorful, grain-textured collage of various items scattered across a light blue background. The items include a green bucket hat, a wooden clothes hanger, a blue-and-red striped collar, blue sunglasses, denim shorts, a yellow price tag, a red high-heele

Google AI Blog · 2026-06-03T13:00

模型发布 中性

How we used Gemini to build Google I/O 2026

A collage of I/O-related images, including the Antigravity Coffee Co. pop-up, a colorful jellyfish and a still from the Timmy TPU video. The word AI repeats three times on the left of the image, and there are colorful icons, including a sparkle, as well.

Google AI Blog · 2026-06-01T16:00

偏多

Take our I/O 2026 quiz, vibe coded in Google AI Studio.

We used Google AI Studio to vibe code a quiz about our top I/O 2026 announcements.

Google AI Blog · 2026-05-29T19:00

模型发布 中性

9 demos of Gemini Omni and Gemini 3.5 in action

Gemini Omni & Gemini 3.5 hero

Google AI Blog · 2026-05-29T17:30

中性

Check out real-life AI prototypes from the Futures Lab.

University of Waterloo students develop AI prototypes like sign language tutors to reshape the future of education and work.

Google AI Blog · 2026-05-29T12:00

中性

Catch up on the Dialogues stage at Google I/O 2026.

Alphabet CEO Sundar Pichai in conversation on the I/O 2026 Dialogues stage

Google AI Blog · 2026-05-22T18:00

模型发布 偏多

100 things we announced at I/O 2026

Image with the words "Ready, Set, I/O" and a colorful Gemini logo

Google AI Blog · 2026-05-20T19:30

中性

A new experiment brings better group meetings to Google Beam

A small group meeting featuring two in-room participants and three participants on-screen.

Google AI Blog · 2026-05-20T16:45

偏多

Google just redesigned the search box for the first time in 25 years — here's why it matters more than you think.

For a quarter century, the Google search box has been one of the most recognizable interfaces in computing: a thin white rectangle, a blinking cursor, a few typed words, and a list of blue links. On Tuesday, Google will formally retire that paradigm. At its annual I/O developer conference, Google announced a sweeping redesign of the search box itself — the literal text field where billions of queries begin every day — transforming it from a simple keyword input into a dynamic, AI-driven conversat

VentureBeat · 2026-05-19T17:45

偏空

How AI Mode is changing the way people search in the U.S.

A graphic features the text "How people are using AI Mode in the U.S." surrounded by colorful, stylized illustrations of a pencil, planet, banana, gift box, cursor, gamepad, and lipstick on a light blue background.

Google AI Blog · 2026-05-19T17:45

中性

New ways to create and get things done in Google Workspace

Google Workspace

Google AI Blog · 2026-05-19T17:45

模型发布 中性

I/O 2026: Welcome to the agentic Gemini era

a series of colorful icons representing Google I/O against a black background

Google AI Blog · 2026-05-19T17:45

模型发布 中性

Gemini 3.5: frontier intelligence with action

Gemini 3.5 text and multi-colored star icon on an abstract blue background.

Google AI Blog · 2026-05-19T17:45

中性

A new era for AI Search

Text that reads "The best of a search engine with the best of AI"

Google AI Blog · 2026-05-19T17:45

中性

Everything new in our Google AI subscriptions, fresh from I/O 2026

The image shows colorful icons. The copy reads: More magic in one plan. More power, more perks, same price.

Google AI Blog · 2026-05-19T17:45

偏多

I/O 2026

At Google I/O 2026, we shared how we're making AI more helpful for everyone. See everything we announced.

Google AI Blog · 2026-05-19T17:45

中性

The new AI-powered Google Finance is expanding to Europe.

A screenshot of the AI-powered experience on Google Finance.

Google AI Blog · 2026-05-11T06:00

投融资

基础设施

利好

Railway secures \$100 million to challenge AWS with AI-native cloud infrastructure

Railway, a San Francisco-based cloud platform that has quietly amassed two million developers without spending a dollar on marketing, announced Thursday that it raised \$100 million in a Series B funding round, as surging demand for artificial intelligence applications exposes the limitations of legacy cloud infrastructure. TQ Ventures led the round, with participation from FPV Ventures, Redpoint, and Unusual Ventures. The investment values Railway as one of the most significant infrastructure sta

VentureBeat · 2026-01-22T14:00

模型发布

AI Agent

开源/评测

偏多

Claude Code costs up to \$200 a month. Goose does the same thing for free.

The artificial intelligence coding revolution comes with a catch: it's expensive. Claude Code, Anthropic's terminal-based AI agent that can write, debug, and deploy code autonomously, has captured the imagination of software developers worldwide. But its pricing — ranging from \$20 to \$200 per month depending on usage — has sparked a growing rebellion among the very programmers it aims to serve. Now, a free alternative is gaining traction. Goose, an open-source AI agent developed by Block

VentureBeat · 2026-01-19T14:00

监管政策

偏多

Listen Labs raises \$69M after viral billboard hiring stunt to scale AI customer interviews

Alfred Wahlforss was running out of options. His startup, Listen Labs, needed to hire over 100 engineers, but competing against Mark Zuckerberg's \$100 million offers seemed impossible. So he spent \$5,000 — a fifth of his marketing budget — on a billboard in San Francisco displaying what looked like gibberish: five strings of random numbers. The numbers were actually AI tokens. Decoded, they led to a coding challenge: build an algorithm to act as a digital bouncer at Berghain, the Berlin high

VentureBeat · 2026-01-16T14:01

AI Agent

偏多

Salesforce rolls out new Slackbot AI agent as it battles Microsoft and Google in workplace AI

Salesforce on Tuesday launched an entirely rebuilt version of Slackbot, the company's workplace assistant, transforming it from a simple notification tool into what executives describe as a fully powered AI agent capable of searching enterprise data, drafting documents, and taking action on behalf of employees. The new Slackbot, now generally available to Business+ and Enterprise+ customers, is Salesforce's most aggressive move yet to position Slack at the center of the emerging "a

VentureBeat · 2026-01-13T13:00

模型发布

AI Agent

偏多

Anthropic launches Cowork, a Claude Desktop agent that works in your files — no coding required

Anthropic released Cowork on Monday, a new AI agent capability that extends the power of its wildly successful Claude Code tool to non-technical users — and according to company insiders, the team built the entire feature in approximately a week and a half, largely using Claude Code itself. The launch marks a major inflection point in the race to deliver practical AI agents to mainstream users, positioning Anthropic to compete not just with OpenAI and Google in conversational AI, but with Microso

VentureBeat · 2026-01-12T11:30

模型发布 开源/评测 偏多

Nous Research's NousCoder-14B is an open-source coding model landing right in the Claude Code moment

Nous Research, the open-source artificial intelligence startup backed by crypto venture firm Paradigm, released a new competitive programming model on Monday that it says matches or exceeds several larger proprietary systems — trained in just four days using 48 of Nvidia's latest B200 graphics processors. The model, called NousCoder-14B, is another entry in a crowded field of AI coding assistants, but arrives at a particularly charged moment: Claude Code, the agentic programming tool from ri

VentureBeat · 2026-01-07T20:00

AI Agent 中性

Which Agent Causes Task Failures and When? Researchers from PSU and Duke explores automated failure attribution of LLM Multi-Agent Systems

In recent years, LLM Multi-Agent systems have garnered widespread attention for their collaborative approach to solving complex problems. However, it's a common scenario for these systems to fail at a task despite a flurry of activity. The post Which Agent Causes Task Failures and When? Researchers from PSU and Duke explores automated failure attribution of LLM Multi-Agent Systems first appeared on Synced.

SyncedAI · 2025-08-14T06:31

偏多

MIT Researchers Unveil “SEAL” : A New Step Towards Self-Improving AI

MIT introduces SEAL, a framework enabling large language models to self-edit and update their weights via reinforcement learning. The post MIT Researchers Unveil “SEAL” : A New Step Towards Self-Improving AI first appeared on Synced.

SyncedAI · 2025-06-16T12:58

AI Agent 中性

Researchers from PSU and Duke introduce “Multi-Agent Systems Automated Failure Attribution

"Automated failure attribution" is a crucial component in the development lifecycle of Multi-Agent systems. It has the potential to transform the challenge of identifying "what went wrong and who is to blame" from a perplexing mystery into a quantifiable and analyzable problem. The post Researchers from PSU and Duke introduce “Multi-Agent Systems Automated Failure Attribution first appeared on Synced.

SyncedAI · 2025-06-16T07:39

模型发布 利好

DeepSeek-V3 New Paper is coming! Unveiling the Secrets of Low-Cost Large Model Training through Hardware-Aware Co-design

A newly released 14-page technical paper from the team behind DeepSeek-V3, with DeepSeek CEO Wenfeng Liang as a co-author, sheds light on the “Scaling Challenges and Reflections on Hardware for AI Architectures.” The post DeepSeek-V3 New Paper is coming! Unveiling the Secrets of Low-Cost Large Model Training through Hardware-Aware Co-design first appeared on Synced.

SyncedAI · 2025-05-15T17:58

模型发布 开源/评测 利好

DeepSeek Unveils DeepSeek-Prover-V2: Advancing Neural Theorem Proving with Recursive Proof Search and a New Benchmark

DeepSeek AI releases DeepSeek-Prover-V2, an open-source LLM for Lean 4 theorem proving. It uses recursive proof search with DeepSeek-V3 for training data and reinforcement learning, achieving top results on MiniF2F. The post DeepSeek Unveils DeepSeek-Prover-V2: Advancing Neural Theorem Proving with Recursive Proof Search and a New Benchmark first appeared on Synced.

SyncedAI · 2025-04-30T15:46

模型发布 利好

Can GRPO be 10x Efficient? Kwai AI's SRPO Suggests Yes with SRPO

Kwai AI's SRPO framework slashes LLM RL post-training steps by 90% while matching DeepSeek-R1 performance in math and code. This two-stage RL approach with history resampling overcomes GRPO limitations. The post Can GRPO be 10x Efficient? Kwai AI's SRPO Suggests Yes with SRPO first appeared on Synced.

SyncedAI · 2025-04-24T02:30

投融资 开源/评测 利好

Zhipu.AI's Open-Source Power Play: Blazing-Fast GLM Models & Global Expansion Ahead of Potential IPO

Zhipu.AI open-sources faster GLM models (8x speedup), launches Z.ai, aiming for global expansion, potentially ahead of IPO. The post Zhipu.AI's Open-Source Power Play: Blazing-Fast GLM Models & Global Expansion Ahead of Potential IPO first appeared on Synced.

模型发布 利好

DeepSeek Signals Next-Gen R2 Model, Unveils Novel Approach to Scaling Inference with SPCT

DeepSeek AI, a prominent player in the large language model arena, has recently published a research paper detailing a new technique aimed at enhancing the scalability of general reward models (GRMs) during the inference phase. The post DeepSeek Signals Next-Gen R2 Model, Unveils Novel Approach to Scaling Inference with SPCT first appeared on Synced.

SyncedAI · 2025-04-11T14:43

趋势信号

[强烈看多] AI Agent 从概念走向企业规模化部署

Gartner预测40%企业应用将嵌入Agent

[看多] 开源模型性价比碾压闭源

DeepSeek V4成本仅为GPT-5的1/30

[中性] 企业端爆发 vs 个人端涨价

行业从流量争夺转向价值收割

[风险] AI电力危机

Gartner预计2027年40%受电力限制

市场传导链

AI 监管加强
政策层



合规成本上升
产业层



安全厂商受益
资本层

主流模型价格对比

厂商	模型	输入	输出	上下文	备注
xAI	Grok 4.1	\$0.20	\$0.50	-	最便宜前沿模型
DeepSeek	V3.2	\$0.27	\$1.10	1M+	性价比之王
OpenAI	o4-mini	\$1.10	\$4.40	-	
Google	Gemini 3.1 Pro	~\$1.25	~\$10.00	2M	推理翻倍
OpenAI	GPT-5	\$1.25	\$10.00	400K	
Anthropic	Sonnet 4.6	\$3.00	\$15.00	1M	中端逆袭
Anthropic	Opus 4.6	\$15.00	\$75.00	200K	最贵旗舰

2026-06-15T03:56

行业首款开源鸿蒙消费级人形机器人：松延动力 N2 亮相华为开发者大会 HDC 2026

IT之家 6 月 15 日消息，行业首款开源鸿蒙消费级人形机器人——松延动力的 OpenHarmony 版 N2 亮相上周末的华为开发者大会 HDC 2026。松延动力表示，当前消费级人形机器人普遍面临应用生态与长期发展的多重痛点。多数产品仍停留在独立硬件终端的形态，难以与手机、智慧屏、智能家居、教育硬件等主流智能设备实现深度互联互通，久而久之形成生态孤岛，极大限制了多场景融合体验的打造。与此同时，传统系统架构下的跨设备开发流程繁琐，网络协议适配、异构数据转换等工作大幅抬高研发门槛与成本，拉长产品迭代周期。再加之外置系统普遍体积偏大、资源占用率高，与人形机器人轻量化、低功耗的行业发展趋势相悖，这些问题都成为制约消费级人形机器人规模化普及、业态升级的关键阻碍。松延动力产品副总裁唐子洋认为，依托开源鸿蒙微内核架构、高安全性与跨终端联动优势，AI 系统智能体语音交互控制、长程任务自规划能力，机器人未来可与家庭场景下的各类智能终端无缝协作，既能化身家庭智能管家，联动家电、安防设备打造全屋智能体验，也能依托多设备协同能力，强化教育陪伴、互动教学等核心场景体验。IT之家从大会获悉，松延动力还宣

2026-06-15T03:26

东风纳米 01 智趣版纯电小车 6 月 16 日上市，CLTC 纯电续航 305km

IT之家 6 月 15 日消息，东风汽车宣布旗下纳米 01 纯电小车智趣版将于 6 月 16 日（明天）正式上市，该车 CLTC 续航里程为 305km，主要在内饰方面同步新近上市的纳米 01 Cross 设计，换用了长条形的全液晶仪表盘以及尺寸更大的中控大屏。外观方面，该车前脸配备分布式灯组，拥有贯穿式日间灯；车身使用隐藏式门把手和低风阻轮圈；汽车尾灯棱角分明，内置圆形 LED 光源，高位扰流板经过上扬处理，不乏运动感。该车座舱采用了纳米 01 Cross 同款设计，换用 15.6 英寸 2.5K 高色域超清护眼智能屏、8.88 英寸液晶组合仪表，有望支持 540° 全景鸟瞰、全场景泊车与循迹倒车等 L2 级组合辅助驾驶功能。规格方面，该车长宽高分别为 4020x1810x1570mm，轴距为 2663mm。不过官方尚未公布新款车型的动力信息，作为参考，现款在售车型采用的单电机前驱布局，电机最大功率为 70kW，峰值扭矩为 160N·m。悬架系统采用了该级别车型中常见的前麦弗逊后扭力梁的组合。

2026-06-15T02:57

小米米家扫拖机器人 6 Pro 国补探新低：水箱版 2192 元、上下水版 2518 元

小米米家扫拖机器人 6 Pro 开售于 2026 年 4 月 8 日，水箱版定价 3799 元、上下水版定价 4199 元。京东 618 大促期间，按下方步骤下单，叠加 PLUS 立减 + 15% 国补后，水箱版 2192.34 元、上下水版 2518.72 元：满 1 元立打 8.2 折券：点此查看。加入购物车时选择“以旧换新”补贴 200 元（可选电动牙刷），京东回收不看旧物品相，老、旧、坏电器皆可回收。晒单再返 20 元京豆。京东小米 米家扫拖机器人 6 Pro 水箱版券后 2192.34 元领 8.7 折券京东小米 米家扫拖机器人 6 Pro 上下水版券后 2518.72 元领 8.7 折券京东无门槛红包至高 26618 元，支持即领即用：点此抽取。IT之家从商品页面获悉，这款新品支持 AI 三摄全景识别，可识别数据线、拖鞋、沙发等 280 种常见障碍物和家具，最细可识别 3mm 耳机线；自研 AI 脏污识别算法，高效识别猫砂、狗粮、牛奶等 47 种干湿脏污，并智能匹配清洁策略；配备高清 AI 摄像头，可用米家 App 在视频模式下，随时查看宠物状态。新品支持小米澎

2026-06-15T02:56

山灵 Moria「生命之树」钛金振膜单动圈耳机本周上市：搭复合振膜单元、钛冰晶工艺面板，4998 元

IT之家 6 月 15 日消息，山灵现已公布一款 Moria「生命之树」钛金振膜单动圈耳机，将于本周上市，定价为 4998 元。该耳机面板采用独特的“钛冰晶”工艺，每片纹理均独一无二，视觉效果如极夜星河或冰面裂纹，具有一定艺术感。同时耳机腔体由高精度 CNC 纯钛打造，轻盈坚固。其标配 5N 长结晶无氧铜镀银线，支持 0.78mm 2PIN 接口，方便用户更换升级线材。官方强调，该耳机搭载“超级动圈”单元，核心为一枚经过 60 层溅射工艺打造的钛金复合振膜，结合 PEK 基材与物理真空直流磁控溅镀技术，在保持柔顺性的同时提升刚性，显著增强解析力。同时，该耳机采用三内腔双磁路动圈架构，通过两个独立后腔室优化气流控制，配合双钕磁铁系统，号称能够实现均衡三频响应与自然的声动态。

2026-06-15T02:48

OPPO ColorOS 16 六月更新公布：锁屏岛新增酷狗音乐等，6 月 30 日前完成推送

IT之家 6 月 15 日消息，OPPO 今日公布 ColorOS 16 六月更新内容：AI 流体云再上新，锁屏岛新增酷狗音乐，流体云新增超级课程表、日历课程提醒；小布建议足球观赛卡，小布助手观赛智能体；还有飞牛私有云无感连接、OPPO Bubble 潮玩自拍屏、全新端午节水印等更多功能。IT之家整理如下：AI 流体云新增支持在小布建议和流体云上，显示超级课程表应用的课程信息支持在小布建议和流体云上，显示支付宝停车服务的计费变动支持在流体云和锁屏岛上，显示酷狗音乐歌词AI 功能新增小布建议足球观赛卡支持通过内容传送门，将图片、文本、文档拖拽存入小布记忆小布助手接入京东、京东金融、飞常准、移动咪咕智能体服务，支持展示对应服务卡片，点击即可查看详情无需跳转第三方应用支持在小布建议上，显示天气预警服务小布建议支持在桌面添加 AI 简报卡片优化小布记忆旅行合集功能，旅行地图支持依据行程生成路线规划图，并在地图中同步呈现景点周边美食小布助手的圈选体验，支持调整圈选框大小与位置，升级圈选光效通信与互联新增设备空间支持添加和管理飞牛私有云设备（NAS）相册支持与飞牛私有云（NAS）无感连接，并可在系

2026-06-15T02:46

长城哈弗 H6 经典版 2026 款上市，限时优惠换新价 8.19 万元起

IT之家 6 月 15 日消息，长城哈弗 H6 经典版 2026 款今日正式上市，上市价 9.99 万元，限时优惠换新价 8.19 万元起。该车限时可享 8000 元现金礼、最高 1 万元置换补贴，以及首任个人非营运车主发动机、变速器的核心零部件终身质保。长城哈弗 H6 经典版 2026 款新增“远山灰”外观配色，尺寸为 4653*1886*1730mm，轴距 2738mm；配有 1.5T 发动机 + 7DCT 湿式双离合变速器，最大功率 135kW，最大扭矩 275N·m。该车配有 14.6 英寸中控屏、手车互联，中控 UI 升级、高德地图红绿灯倒计时；支持 ACC 自适应巡航、540° 全景影像。IT之家查询配置表获悉，长城哈弗 H6 经典版 2026 款搭载前麦弗逊式独立悬架 / 后多连杆式独立悬架，标配 6 安全气囊，车身高强度钢占比约 72%。该车还新增电尾门、前排座椅电动调节、前排座椅加热，配有可开启全景天窗、8 扬声器、外后视镜加热以及自动折叠。

2026-06-15T01:04

Claude Code Guide 2026: 25 Features with Examples + Demo

Claude Code is a layered agentic coding tool, not a single chat prompt. This guide breaks down 25 features, from CLAUDE.md, skills, subagents, and hooks to MCP and Auto Mode. It includes a comparison table, working code examples, real use cases, and an interactive demo you can try. The post Claude Code Guide 2026: 25 Features with Examples + Demo appeared first on MarkTechPost.

2026-06-15T00:00

A Deep Reinforcement Learning (DRL)-Based Transformer Method for Solving the Open Shop Scheduling Problem

arXiv:2606.13682v1 Announce Type: new Abstract: The open shop scheduling problem (OSSP) arises in many industrial and service settings but remains computationally challenging as the number of jobs and machines increases. While exact methods quickly become intractable, classical dispatching rules and metaheuristics may require substantial tuning to maintain solution quality at large scales. This study develops a Transformer-based scheduling policy for OSSP using an encoder-decoder architecture w

2026-06-15T00:00

Orchestra-o1: Omnimodal Agent Orchestration

arXiv:2606.13707v1 Announce Type: new Abstract: The recent success of agent swarms has shifted the paradigm of large language model (LLM)-based agents from single-agent workflows to multi-agent systems, highlighting the importance of agent orchestration for task decomposition and collaboration. However, existing orchestration frameworks are limited to a narrow set of modalities and struggle to generalize to more complex settings where heterogeneous modalities coexist and interact. This limitati

2026-06-15T00:00

Hybrid Open-Ended Tri-Evolution Makes Better Deep Researcher

arXiv:2606.13710v1 Announce Type: new Abstract: Deep research and agent evolution serve as de-facto tasks for AI agents in real-world applications toward artificial general intelligence. The former enables autonomous retrieval and integration of information in open-ended environments to tackle open-ended research tasks, yet it is constrained by the static parametric deep research capabilities of agent systems. The latter allows agents to autonomously interact with the environment to gain experi

2026-06-15T00:00

WorkBench Revisited: Workplace Agents Two Years On

arXiv:2606.13715v1 Announce Type: new Abstract: The best agent on WorkBench in March 2024, GPT-4, completed 43% of tasks and took an unintended harmful action, such as emailing the wrong person, on 26% of them. We re-visit the benchmark in June 2026 and find that the best agent to date, Claude Opus 4.8, completes 89% and takes an unintended harmful action on 2.5%. Aside from this considerable progress in frontier agent performance, three things stand out. First, capability and safety go together

2026-06-15T00:00

Refusal Beyond a Single Direction: A Preliminary Comparison of Diff-in-Means and INLP

arXiv:2606.13720v1 Announce Type: new Abstract: Ardit et al. (2024) has shown that refusal in safety fine-tuned chat models is mediated by a single linear direction in the residual stream, recoverable by a difference-in-means (DiM) of harmful and harmless activations. We compare DiM-based interventions (activation addition and directional ablation) with two interventions derived from Iterative Nullspace Projection (INLP) -- nullspace projection and counterfactual flipping -- on five open-weigh

2026-06-15T00:00

YeasierAgent: Agentic Social Sandbox as a Canvas for Intent-Driven Creation of Platform-Agnostic Symbiotic Agent-Native Applications

arXiv:2606.13722v1 Announce Type: new Abstract: This paper introduces YeasierAgent, an application-building paradigm based on symbiotic agents, narrative worlds, and scene-aware interaction. It challenges the conventional device-coupled model of software by redefining applications as collaborative spaces among users, agents, and worlds. We present a system architecture that achieves two primary contributions: (1) enabling the rapid, cross-platform construction of agent-native applications by ut

2026-06-15T00:00

TwinBI: An Agentic Digital Twin for Efficient Augmented Interactions with Business Intelligence Dashboards

arXiv:2606.13731v1 Announce Type: new Abstract: Business intelligence (BI) increasingly combines dashboard interaction with LLM-based assistance, but these two modes often fall out of sync during multi-step analysis. As users switch between direct dashboard manipulation and natural-language queries, it becomes difficult to preserve a consistent analytical state across filters, hierarchies, metrics, and chart context. We present TwinBI, an agentic digital-twin framework that couples an LLM-based

2026-06-15T00:00

AI Receptivity or AI Adoption Breadth? A Tool-Specific Reanalysis of the Lower-Literacy/Higher-Usage Link

arXiv:2606.13734v1 Announce Type: new Abstract: Recent evidence reported by Tully, Longoni, and Appel (2025) suggests that lower artificial intelligence (AI) literacy predicts greater receptivity toward AI. We revisit this claim using the public data from Study 3 of that article, which measures past usage of five AI tool categories on a five-point frequency scale. We first reproduce the negative association between AI literacy and aggregate AI usage using OLS on participant-level averages, bina